

Predicting Pipeline Failures

In collaboration with Veolia Research & Innovation

Thomas Bordier, Christophe Lanternier

Spring 2017

Table of contents

1. Data presentation
2. How to work with an unbalanced data set?
3. Data Representation
4. Classification Algorithms used
5. Thank you for attention, any questions?

Data presentation

Data presentation

	Feat1	Feat2	Feat3	Feat4	Length	Year Constr.	Year Last Fail.
1	T	IAB	-0.209841	C	3.5972	2001	NaN
2	T	U	2.184992	M	4.0547	1965	NaN

Table 1: Sample from the raw dataset

- Low dimensionality
- No obvious invariance
- Highly unbalanced dataset: **1/200** pipeline fails.

Basic Data Observations




Figure 1: Feature 1 observations

Basic Data Observations




Figure 2: Feature 2 observations

Basic Data Observations




Figure 3: Feature 3 observations

Basic Data Observations






Figure 4: Feature 4 observations

Basic Data Observations



(a) Broken Pipes



(b) All Pipes

Figure 5: Pipe Length

How to work with an unbalanced data set?

There are two main issues to be tackled:

- Re-establish balance in the dataset
- Detect and prevent over-fitting

detecting and preventing over-fitting

Detecting over-fitting:

- Implementation of a randomized train/test split
- Observation of mean scores and variance over 20 experiences

Preventing over-fitting:

- Focus on regularized models (Adaboost with low number of estimators, Logistic Regression with high regularization)

These precautions allowed a big jump in performance and ranking.

Reestablishing balance

We considered two ways of doing so:

- Simple duplication
- Synthetic Minority Over-sampling TEchnique [1]

SMOTE

Intuitive comprehension of the SMOTE algorithm:




Figure 6: SMOTE explained

Performance comparison




Figure 7: Mean score for 3, 5, 7, 10 and 20 estimators on Adaboost

Performance comparison

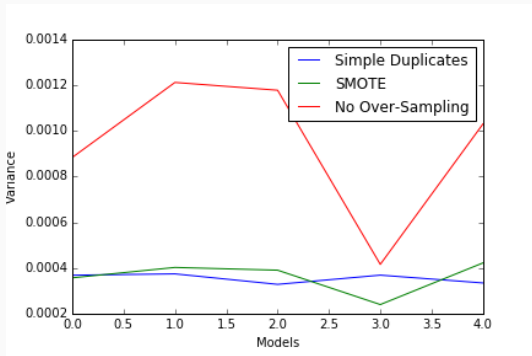


Figure 8: Variance for 3, 5, 7, 10 and 20 estimators on Adaboost

Data Representation

How to go beyond basic features?

	F3	Length	Year Const.	Year Last Obs.	P	T	IAB	O	U	C	D	Dr	M
1	-0.001506	0.000494	0.003946	-0.007229	0	1	1	0	0	1	0	0	0
2	0.015677	0.000557	0.014094	-0.007229	0	1	0	0	1	0	0	0	1

Table 2: Basic preprocessing

- **Idea 1:** non-linear combinations
- **Idea 2:** neural network embeddings

Non-linear combinations

Implement binary operations such as:

- AND
- OR

For pairs and triplets for categorical variables

	F3	Length	Year Const.	Year Last Obs.	P	T	IAB	O	U	C	D	Dr	M
1	-0.001506	0.000494	0.003946	-0.007229	0	1	1	0	0	1	0	0	0
2	0.015677	0.000557	0.014094	-0.007229	0	1	0	0	1	0	0	0	1

Table 3: Basic preprocessing

Non-linear combinations

Metrics	Raw data	Pair 'and'	Pair 'or'	Pair 'and' + 'or'	Triple 'and' + 'or'
Mean AUC Score	0.840	0.860	0.857	0.860	0.865
Variance AUC Score	4.38×10^{-4}	4.33×10^{-4}	7.95×10^{-4}	4.21×10^{-4}	6.50×10^{-4}

Table 4: Evolution of our estimated score with binary operations, with AdaBoost

Careful

Curse of dimensionality when too many combinations (13 features to 900 features approx.)

Neural Network feature

- We gave neural network generated features, a try, without much success.




Figure 9: Evolution of accuracy with epochs

Classification Algorithms used

Classification algorithms

- Logistic Regression
- Decision Trees
- SVM
- Neural Networks
- **AdaBoost**

Focus on best performing algorithms

- Logistic Regression
- AdaBoost

Logistic Regression




Figure 10: Evolution of score with penalization

Logistic Regression with L_1 penalty

Feat	F3	Length	Years Last. Fail.	TC and	TDr or	IABM or	ODr or	OM or	POC or	TIABDr or	TODr or	TUD and
Importance	-69	125	402	-1	-1	-6	-1	-6	-1	-1	-1	-1

Table 5: Feature importance in model

Interpretability

L_1 penalty leads to sparsity

This method leads to a score around 0.76 locally, 0.74 on the website.

- Model update:

$$C_m(x_i) = C_{(m-1)}(x_i) + \alpha_m k_m(x_i)$$

- Weight update and loss function:

$$E = \sum_{i=1}^N w_i^{(m)} e^{-y_i \alpha_m k_m(x_i)} \text{ with } w_i^{(m)} = e^{-y_i C_{m-1}(x_i)}$$






Our performance depended on:

- Effective **over-fitting control**
- Effective data **augmentation**
- Effective data **representation**
- **Optimization** of hyperparameters for AdaBoost algorithm

Final Performance

Ranked 2nd - Final Score : 0.8845

Thank you for attention, any
questions?

-  Nitesh V. Chawla, Kevin W. Bowyer, Lawrence O. Hall, W. Philip Kegelmeyer. SMOTE: Synthetic Minority Over-sampling Technique. *Journal of Artificial Intelligence Research* 16, 321–357, 2002.
-  Yoav Freund and Robert E. Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences* , 55(1):119-139, August 1997.
-  Robert E. Schapire. *Explaining Adaboost*
-  Gareth James, Daniela Witten, Trevor Hastie, Robert Tibshirani, *An Introduction to Statistical Learning*, 2013.
-  Jerome H. Friedman, Greedy Function Approximation: A Gradient Boosting Machine. *IMS 1999 Reitz Lectures*. 1999